

Computer Networking

Assignment 1: Presentation

HTTP and Web Servers

Due date: **Friday, April 9, 2021 at 17:00**

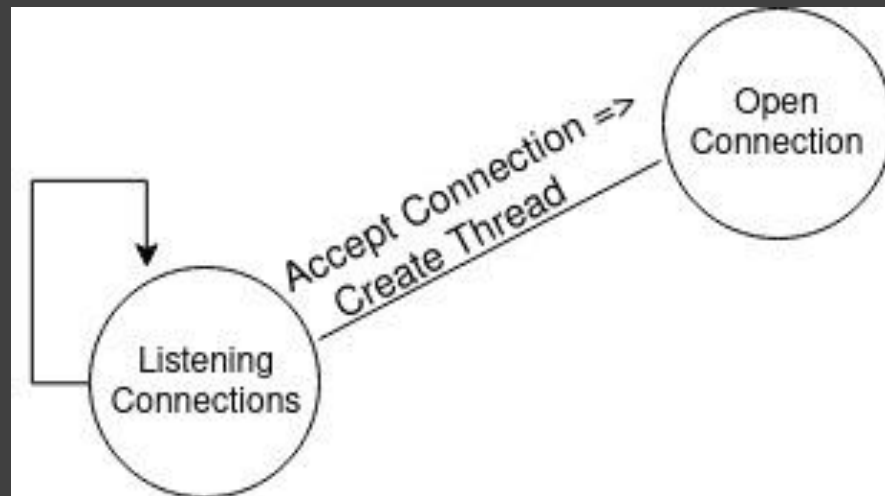
USI Lugano, March 24, 2021

1. Introduction

- The objective of this assignment is to develop a Web server.
- A Web server is a program that serves (provides) Web resources, and satisfies the HTTP requests of clients that want to access those resources.
- The assignment will cover the expected functionality from the server, as well as some HTTP and Web server concepts that are useful to know.

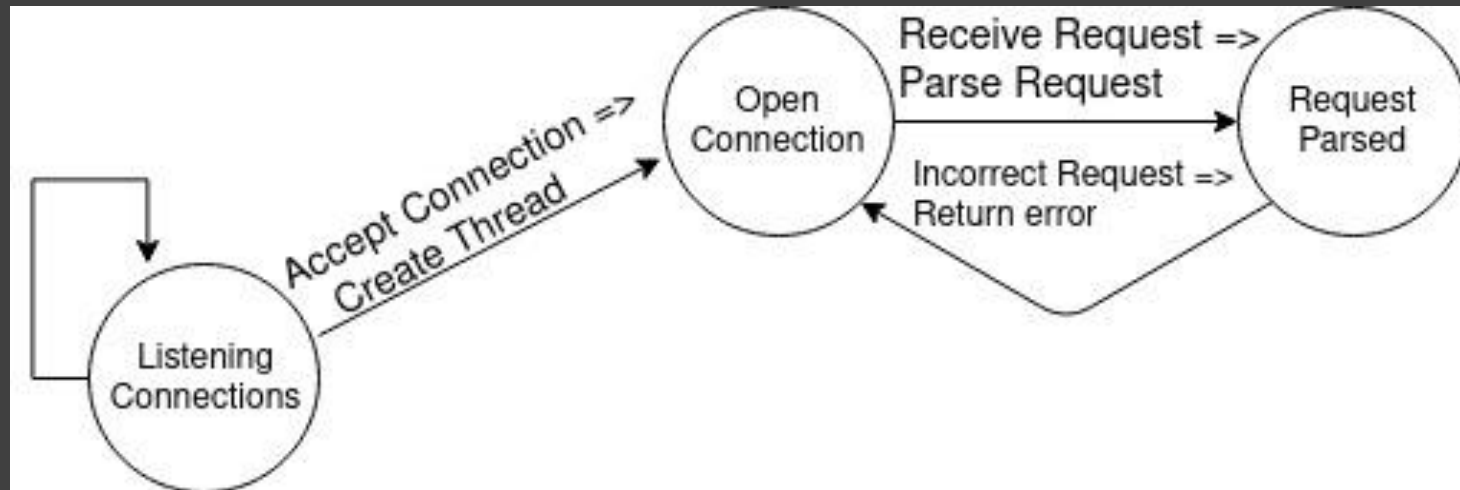
1.1 Server execution flow (1)

- The Web server should behave as follows. A main thread of the server will create a socket and will be infinitely looping, listening for TCP connections. When a new connection arrives, the server accepts it and creates a new thread (this is an optional task) that will be in charge of handling the connection.



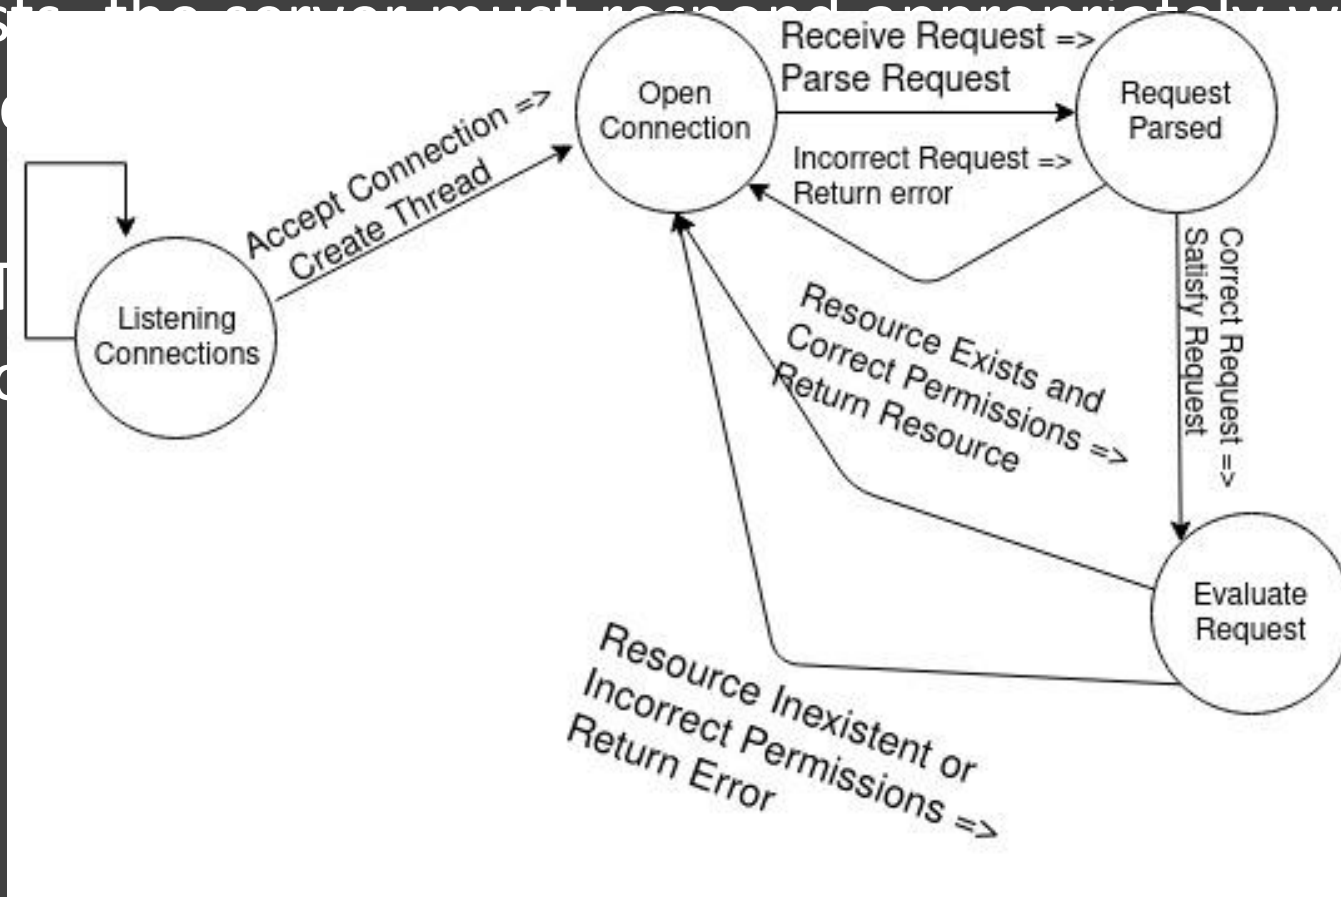
1.1 Server execution flow (2)

- The new thread will start receiving the HTTP requests through the opened TCP connection. The server must check the correctness of each request and return an error (which?) if one of them is malformed (missing mandatory headers, unknown methods, etc.). If the request is correct the server must continue and try to satisfy it.



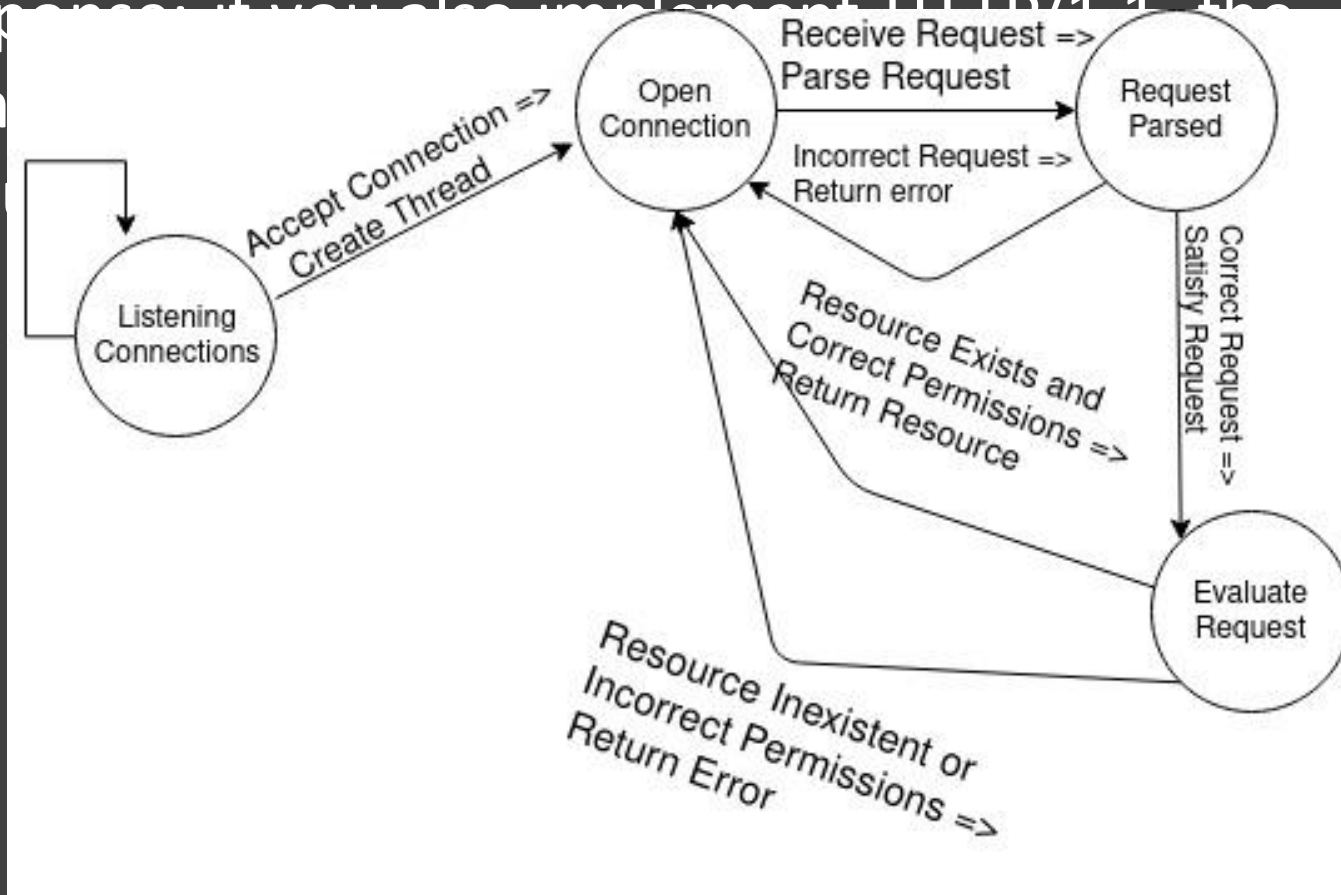
1.1 Server execution flow (3)

- If the resource specified in the request does not exist, a new error will be returned (which?); if the resource exists, the server must respond appropriately with the function in the GET request or a resource



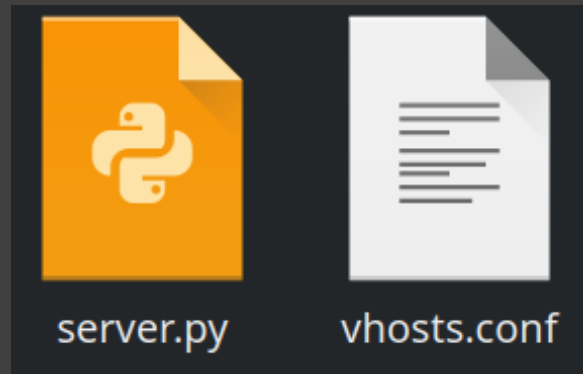
1.1 Server execution flow (4)

- If you only do the HTTP/1.0 implementation, the server must close the connection after giving a response; if you also implement HTTP/1.1, the connection can be kept open for subsequent requests.



2. Tasks – 2.1 Virtual Hosts (1)

- Although it is easy to imagine sites like Google requiring more than one server to answer all of the world's queries, the converse is also true: sometimes one server answers requests for more than one site.
- This is called virtual hosting, and the configuration is usually done through a file. We will call this file *vhosts.conf*, and it must be placed in the same folder where the server is executed.

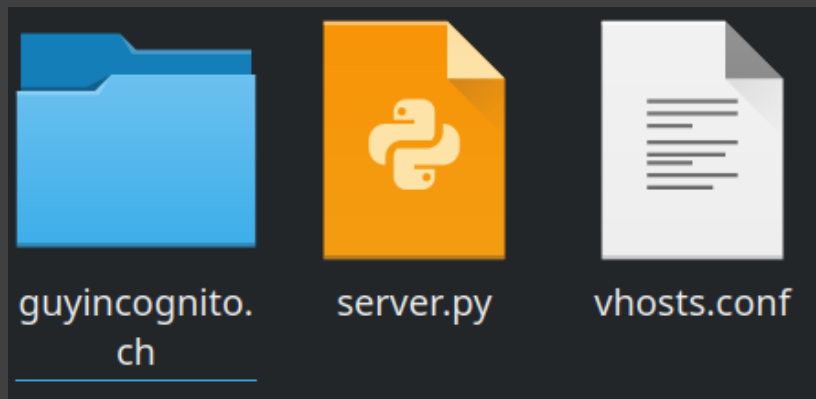


2. Tasks – 2.1 Virtual Hosts (2)

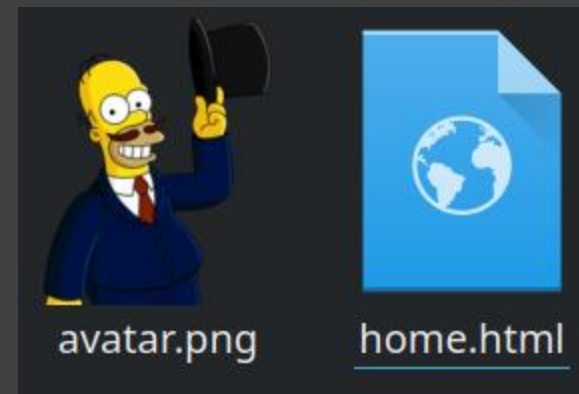
- Your server should host one site per each group member. Each site must be hosted in its own directory, and the name of the directory must match the domain of the site (i.e., if my name were Guy Incognito, and I wanted to access my site at guyincognito.ch, the folder with the site's files should have the same name).
- Each site must contain at least one HTML file and have an image in its content. The HTML file should include a short text (approximately of 400 words) referring to the site owner (the group member), with the image that relates to the text.

2. Tasks – 2.1 Virtual Hosts (3)

- As an example, assume a Web server is created to serve the site *guyincognito.ch*. In the same folder where the server is executed, there must be the *vhosts.conf* file and a folder called *guyincognito.ch*, with the site's files (HTML, picture) inside. Therefore, the file/folder structure (in this example with only one site) should be as follows:



Server Folder



Guyincognito.ch folder

2. Tasks – 2.1 Virtual Hosts (4)

- The configuration of the virtual hosts through the `vhosts.conf` must be done in the following way. Each site must occupy one line and be a list of comma-separated values formatted as follows:

domain,entry point file,member fullname,member email



The entry point of the site in this example should be *home.html*. This means that if a Web client makes a request for the resource `/` at host *guyincognito.ch*, the server must reply with the content of the *home.html* file. The *home.html* file must be placed within the *guyincognito.ch* folder. Therefore, the line of the `vhosts.conf` file for the example site should be:

`guyincognito.ch,home.html,Guy Incognito,guy.incognito@usi.ch`

Since the assignment asks for three sites to be created (one per group member), the corresponding file should have three such lines. If your group has a different size, you must put that many sites.

2. Tasks – 2.1 Virtual Hosts (5) - TASKS

Task A (5 points)

The server should be able to read from a vhosts.conf file the information regarding the sites it serves.

Task B (10 points)

The submission must include one site per group member.

2. Tasks – 2.2 HTTP protocols (1)

- The Web server should handle at least HTTP/1.0 requests.
- Optionally, the server can be made capable of handling also HTTP/1.1 requests.
 - If such a functionality is implemented, up to 10 additional points will be assigned for this task (on top of the maximum points available for the task).
- Although there are several differences between both versions of the protocol, in the context of this assignment we only care about two

2. Tasks – 2.2 HTTP protocols (2)

- HTTP/1.0 expects a new TCP connection for each request; i.e., the TCP connection is automatically closed after receiving an HTTP request and answering it.
 - On the contrary, HTTP/1.1 supports persistent connections; i.e., the connection is kept open and the server receives and answers several requests per connection, until it is explicitly closed.
- HTTP/1.0 has an optional host header.
 - For HTTP/1.1, that header is instead mandatory. This means that for HTTP/1.0, if no host is specified, the server must use as default the first site defined in the vhosts.conf file.

2. Tasks – 2.2 HTTP protocols (3) - TASKS

Optional Task A (10 extra points)

The server should handle HTTP/1.1 connections.

2. Tasks – 2.3 HTTP methods (1)

The Web server should answer to a subset of the available HTTP methods: GET, PUT and DELETE. Additional information about the functionality of the different HTTP methods is available in the corresponding RFCs for HTTP/1.0 and HTTP/1.1.

- HTTP/1.0: RFC1945, Section 8 and Appendix D.1
- HTTP/1.1: RFC7231, Section 4.3

2. Tasks – 2.3 HTTP methods (2)

Additionally, the server should also include the NTW21INFO method. Requests for this method will only have / as the resource, and the server should reply like in the following example.

Server (response)

```
1 HTTP/1.0 200 OK
2 Date: Wed, 24 Mar 2021 09:30:00 GMT
3 Server: NTW21Instructors
4 Content-Length: 98
5 Content-Type: text/plain
6
7 The administrator of guyincognito.ch is Guy Incognito.
8 You can contact him at guy.incognito@usi.ch.
```

Client (request)

```
1 NTW21INFO / HTTP/1.0
2 Host: guyincognito.ch
3
```


2. Tasks – 2.3 HTTP methods (3) - TASKS

Task C (10 points)

The server should be able to answer well-formed GET requests as specified by the corresponding RFCs.

Task D (10 points)

[...] PUT requests [...]

Task E (10 points)

[...] DELETE requests [...]

Task F (5 points)

The server should be able to answer to NTW21INFO requests as specified by the assignment.

2. Tasks – 2.4 Request format (1)

This section covers information needed to solve Tasks C-F. Therefore, no specific tasks are defined, but the information given here must be considered for solving those tasks.

- Normally the client requests will only consist of 2-4 lines, but the PUT method will also include an entity body.
- The first line is always the request-line, whose format is specified in the RFCs: method to be applied, identifier of the resource, and protocol to use.
- The remaining lines include the headers and the optional body. Although the RFCs define several possible headers for client requests, we are only interested in:
 - *Host* header (optional in HTTP/1.0)
 - *Connection: close* header (when ending connection in HTTP/1.1).
 - *Content-type* and *Content-length* headers. Used in the PUT method, along with the entity body (the content of the file being put into the server).

2. Tasks – 2.4 Request format (2)

- The server should not crash when it finds unexpected headers: it should ignore them (otherwise, you will not be able to test it with a Web browser).
- If there is no Content-length header (no message body), the request ends after a CRLF alone in a line.

```
1 GET /home.html HTTP/1.0
2 Host: guyincognito.ch
3
```

- If there is a Content-length header, the end of the request is after reading that amount of bytes in the message body (after the CRLF).

```
1 PUT /new_file.html HTTP/1.0
2 Host: guyincognito.ch
3 Content-type: text/html
4 Content-length: 57
5
<html>
```

In this case, the server creates a new file called new file.html in the guyincognito.ch folder, with the specified HTML content. If the file was already existing, then it replaces its content with the new content from the request.

2. Tasks – 2.5 Response format (1)

This section covers information needed to solve Tasks C-F. Therefore, no specific tasks are defined, but the information given here must be considered for solving those tasks.

- The responses from the server should always include the *status-line*.
- The *message body* at the end of a reply should only be answered to a GET request.
- Regarding *response-header/header fields*, we only care about the following:
 - *Date* (for GET and DELETE)
 - *Content-length* (for GET, NTW21INFO)
 - *Content-type* (for GET, NTW21INFO)
 - *Server* (responding with the name of your group, for all methods)
 - *Content-location* (for PUT).

2. Tasks – 2.5 Response format (2)

For example, if the client asked for:

```
1 GET /new_file.html HTTP/1.0
2 Host: guyincognito.ch
```

The server would answer with:

```
1 HTTP/1.0 200 OK
2 Date: Wed, 24 Mar 2021 09:45:53 GMT
3 Server: NTW21Instructors
4 Content-Length: 57
5 Content-Type: text/html
6
7 <html>
8 <body>
9 Hello! This is a new file.
10 </body>
11 </html>
```

Note the CRLF character between the *status-line/response-header* and *message body*.

2. Tasks – 2.6 Response codes

It is expected that the Web server correctly handles at least (you are free to cover more :)) the cases that generate the following HTTP status codes: 200, 201, 400, 403, 404, 405.

Task G (10 points)

The server must answer client requests with the correct status code appropriate for the situation. Status codes 200, 201, 400, 403, 404 and 405 must be considered.

2. Tasks – 2.7 Multithreading

- One of the performance metrics used for evaluating Web servers is the number of requests that it can answer per unit of time.
- A server that can only answer to one request at a time will therefore score quite low in this metric.
- To obtain good performance (and extra points!) the server should be multithreaded, i.e., it should be capable of serving multiple requests simultaneously.
- You can find a Java example of a multithreaded server available on iCorsi.

Optional Task B (10 extra points)

The server must be multithreaded.

2. Tasks – 2.8 Extra info

- You can always refer to the HTTP/1.0 ([RFC1945](#)) AND HTTP/1.1 RFCs (in particular, [RFC7230](#) and [RFC7231](#)) for more information about what is expected from a Web server and what it must/should, must not/should not, do.
- We also recommend Mozilla Developer Network documents about [HTTP](#), where you can find information about the [methods](#) , the [response status codes](#), and the [headers](#).
- The information is also given in a more user-friendly way than the RFCs (but, of course, may contain mistakes: the RFCs are the standard).

3. Submission – 3.1 Testing

- The Web server should be able to handle well-formed requests from a Web browser, but also from command-line tools like netcat, telnet or cURL. Make sure the latter work, as they will be used to test your solution.
- In order to test the different sites from your Web browser, you will have to edit the hosts file of your operating system, and point the site domain (e.g., *guyincognito.ch*) to the IP 127.0.0.1 (localhost).
- Then you will be able to input your domain in your Web browser and the request should be answered by your server.

3. Submission – 3.2 Program Execution

- You can find on iCorsi, next to this assignment's PDF, a compressed file with the example of a site's content, along with the vhosts.conf file. If you put this content in the same folder than your server, it should be able to serve the site.
- **Remember that there must be one site per group member.**

3. Submission – 3.3 Available configuration example

- Your program must accept the following command-line arguments:
 - PORT NUMBER: although 80 is the well-known port for HTTP requests, this should be configurable. If no port number is provided as an argument, then use 80.

Task H (5 points)

The server must accept a port number as an argument and it must listen for connections on it. Use 80 as default.

3. Submission – 3.4 Submission Instructions (1)

- You may write your solution in Java or Python (**recommended**). Add comments to your code to explain the code itself.
- Package all the source files plus a README file in a single zip or tar archive.
- **It is explicitly forbidden to include any other files or folders**
 - **i.e. no *.directory*, *__MACOSX*, *thumbs.db*, etc!**
- The README file is a report of your assignment. Use it to add general comments, acknowledge external sources of information (code, suggestions, and comments from other students/groups)
- If your implementation has limitations and errors you are aware of (and were unable to fix), then list them as well.

3. Submission – 3.4 Submission Instructions (2)

- The README file must include a list of the group members with the contributions made by each one.

Guy Incognito	Task A, Task B, Optional Task A
Jane Doe	Task B, Task C, Optional Task B
John Doe	Task B, Task D, Task E

- Make sure that you include all the necessary documentation and components to build and run your solution on a standard installation of a Java or Python environment.
- In particular, make sure your solution works with the most basic command-line tools, outside of any integrated development environment.

3. Submission – 3.4 Submission Instructions (3)

- The README file must include a list of the group members with the contributions made by each one.

Guy Incognito	Task A, Task B, Optional Task A
Jane Doe	Task B, Task C, Optional Task B
John Doe	Task B, Task D, Task E

Task I (20 points)

Your code should be in a working state, without needing modifications from the Instructors, and your submission must comply with the specified guidelines. Points will be deducted if the guidelines are not respected, if the code crashes, has compilation errors or does not execute, or if additional tools are needed for compilation/execution (like maven, gradle, etc.).

Task J (15 points)

Your code must be appropriately documented and your submission must include a README file as a report of your work. The report must include compilation and execution steps, plus a table of contributions.

4. Grading

- The assignment has a total of 100 points. If your submission exceeds 100 points, that excess will be added to your future assignments.

Table 1: Score for tasks

Task	Points
A	5
B	10
C	10
D	10
E	10
F	5
G	10
H	5
I	20
J	15

Table 2: Extra score for optional tasks

Task	Points
A	5
B	10

Report

- Along with your code submission, please add a PDF file (1-2 pages) containing the following information
 - Detailed **list of the sources** used to solve the assignment
 - Including conversations with other groups in the class, blog posts, github code repositories, stackoverflow threads, etc.
 - A **table** showing the tasks carried out by each group member
 - If all group members collaborated equally to all tasks, please do specify that
 - Any other information you believe would be relevant to communicate to us



Using references vs Plagiarism

- You are **ALLOWED** to
 - Consult with other students on the assignment
 - Consult public resources such as Web pages, code repositories, etc.
- You are **REQUIRED** to
 - Cite in a complete and exhaustive manner ALL SOURCES you consulted to solve the assignment
 - Explicitly indicate in the code if any part of the code has been copied from others
- Any, even minor, act of plagiarism will be punished with a **0 points** evaluation of the entire assignment for all group members AND with a **notification to the Dean's office**

Using references vs Plagiarism

plagiarism



Pronunciation  /'pleɪdʒərɪz(ə)m/ 

[See synonyms for plagiarism](#)

Translate [plagiarism](#) into Spanish

NOUN

[mass noun]

The practice of taking someone else's work or ideas and passing them off as one's own.

‘there were accusations of plagiarism’

[+ More example sentences](#)

[+ Synonyms](#)

Group Size

- We request a group size of 3
- Groups of 2 will be asked to find a new member ASAP
- Otherwise, the group will be assigned an additional member by the instructors or may be split and its members assigned to other groups

ANY
QUESTIONS
?